

My experience in data management systems and public engagement activities

Saumya Bhatnagar

March 5, 2020

Why DBMS!

Users

DBA

DB Schema

APP PROGRAMMERS

App Software

END USERS

Query App Interface

DBMS

Query Processor

Query Evaluation Engine (DDL Interpreter, DML Compiler, Application Object Code)

Storage Manager

Buffer Manager, File Manager, Transaction Manager

Database

Data files, Data Dictionaries, Indices

DBMS Types

	SQL	NoSQL
High Level Model	ER Model	
Representational Model	Hierarchical (IMS), Relational (Oracle, DB2, SQL Server), Network (IDMS, IMAGE)	
Low-Level Model		

DB Architectures

- Centralized DBMS Architecture
- Client-Server Architecture
- Distributed Database Architecture

Schema Types

- Internal Schema
- Conceptual Schema
- External Schema

Glossary on Keys

Key Types

- Super Key
- Candidate Key
- Primary Key
- Secondary Key
- Foreign Key
- Composite Key
- Compound Key (Composite key with foreign key)
- Alternate Key
- Sort/Control Key
- Surrogate key

Overview

1 DBMS

- MySql, Oracle, Cassandra, HBase, MongoDB

2 Hadoop

- Hadoop Ecosystem
- External Data Storages
- Query Engines

3 Which Data Storage?

4 SQL

- MySql, Vertica

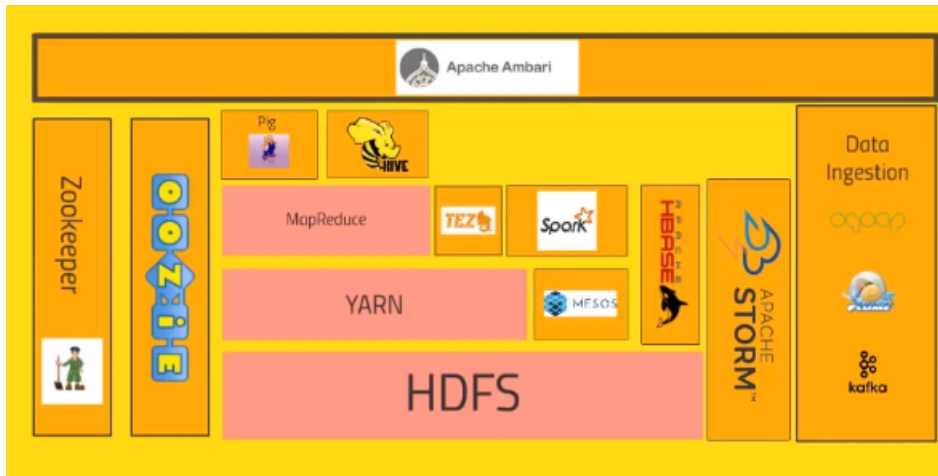
5 NoSQL

- Cassandra with solr
- No one single point of failure

6 APIs

7 Microservices

Hadoop Ecosystem



Query Engines And External data storage

Query Engines



External Data Storage



Clustered Computing Platforms (Mapreduce, Spark)

SPARK

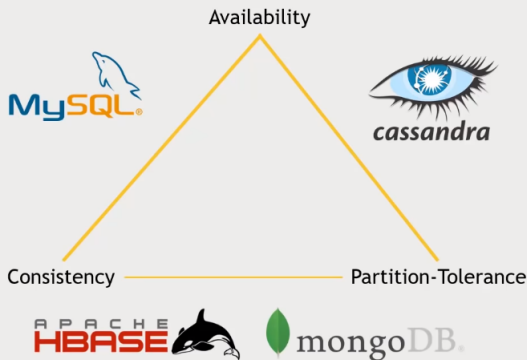
- Distributing queries and trend analysis
- Microbatching for historical analysis
- Loading large datasets into memory
- Running queries against large datasets

Pros & Cons of the databases

Hadoop/Mapreduce	Slow for real time analytics
MongoDB	Global write lock performance concerns
Cassandra (w/o solr)	Query Limitations
Cassandra (w/o solr)	No bother about denormalizing, duplication, access pattern data modelling
Solr	Search capabilities, partial text search, facet queries, geospatial, etc.

Which Data Storage?

CAP considerations



Vertica for Big Data Engineering

Command Type

- | | |
|-------|---|
| ① DDL | ① create, alter, drop, truncate, rename |
| ② DML | ② select, insert, update, delete |
| ③ DCL | ③ grant, revoke |
| ④ TCL | ④ commit, rollback |

Example (Vertica Code Example)

```
SELECT name, class, date,  
RANK() OVER (PARTITION BY class ORDER BY marks desc) AS rank  
FROM student  
WHERE name IS NOT NULL  
AND subject like 'math%'  
AND date > '01/01/2007'  
ORDER BY class;
```

SQL Glossary

- bandwidth=rate of data transfer
- latency=time of data transfer
- 1NF, NF, 3NF, BCNF
- ACID Properties (atomicity, consistency, isolation, durability)
- Lossless Decomposition
- Data Independence
-
-
-
-
-
-

DSE provides integration between Cassandra with Solr

- Storage grid (cassandra) + Search grid(solr)
- Devcenter or cqlsh
- Cassandra cluster handling over 1TB data
- 2 Data Centers
- 3 Servers, with RF of 3
- configure dse.yaml or vassandra.yaml
- Opscenter
- Solr Admin UI gives Solr Index Size
- All Nodes should have solr enabled within DC
- Map collection to dynamic fields
- solr queries have consistency levels

CQL syntax similar to SQL syntax

Example (CQL Code Example)

```
/*create table defining partition, clustering keys*/  
CREATE TABLE student (  
name text, class text, subject text, date timestamp,  
PRIMARY KEY ((name, class), date)  
);
```

Primary key is defined as ((partition keys), clustering/sorting keys)

Example (CQL Code Example)

```
SELECT name, class, date, rank FROM student  
WHERE name IS NOT NULL  
AND subject CONTAINS 'math'  
AND date > '01/01/2007'  
ORDER BY class  
PER PARTITION LIMIT 2;
```

Solr provides full text search, term-search

Example (CQL + Solr Code Example)

```
SELECT name, class, date
FROM student WHERE
solr_query='{ "q": "name:[* TO *] AND subject:math*",
              "fq": "date:[2007-01-01T00:00:00Z TO NOW]",
              "facet": { "field": "class" },
              "sort": "class, marks desc",
              "paging": "driver",
              "timeAllowed": 30000 }'
ALLOW FILTERING;
```

Clustering columns can be defined in WHERE clauses if ALLOW FILTERING is also used even if a secondary index is not created

Cassandra Glossary

- snitch
- Gossip
- Quorum
- num_tokens
- $\text{max_solr_concurrency_per_core} = \text{cpu code} / \text{num solr cores}$
- partitioner
- auto_bootstrap
-
-
-
-
-

SOAP vs REST

Client (Machine Devices - Mobile, desktop) → API Binding → Server
SOAP:

- Stateless
- Slow
- XML

REST:

- Public
- Fast
- Multiple formats

REST:

NODE.js

MongoDB (native js code) - JS based

json format

MongoDB: 2 collection joins, aggregation in mongoDB
instead js for loop can be used

REST vs Bulk

Bulk is built on top of REST
Bulk:

- async
- batches

Email API

Email uses SMTP and Port number

Tight coupling

IOC (Inversion of Control):

Inject email in customer: 1. Property in class 2. Parameter

command

...

read

...

API gateway

Swagger

APIGEE:

- authentication control
- traffic control

Server info... API gateway provides URL

Microservices

Microservices architecture runs on top of STORM/JMS/KAFKA

Storm (handles clustering/distribution)

Kafka (messaging between the grids)

Kafka or Rabbit NQ are message broker URIs

for cache use Redis. Redis is a cache DB

JWT (Json Web token): network calls to DB should be least → Resource Management

YAML → dependent on other services. has details such as name, port, URL, env variables, etc.

Docker - Container

Docker is OS

Containers are VM ware

Cluster has nodes. Nodes has pods. e.g. Pod1, Pod2, Pod3, Pod4 are 4 containers. Pod1 may act as Inst of Service

Dockerfile is image of service and is "Built, deployed and ran" by DevOps

Domain Driven Design

Service bus

Rabbit MQ

Order Service & Domain Service

DDD: Command (message) → Event [Eventual Consistency]

Service bus... message sent to exchange queue via routing key

Pub/Sub Design Pattern

content...

title

content...

Microservices on Docker

content...

Microservices on Kubernetes

content...

Serverless

content...

Thank You!